

# Let's Write A Test!

@joedevivo  
basho technologies

<http://joedevivo.github.io/euc2013>

I love Sweden!



# The Vasa









NORTH

Delta  
Frontier

Jet Blue  
Sun Country





Engineering  
since like forever

Build Thing

Test Thing

Use Thing

# Software Engineering

now we have options!

Test Thing

Build Thing

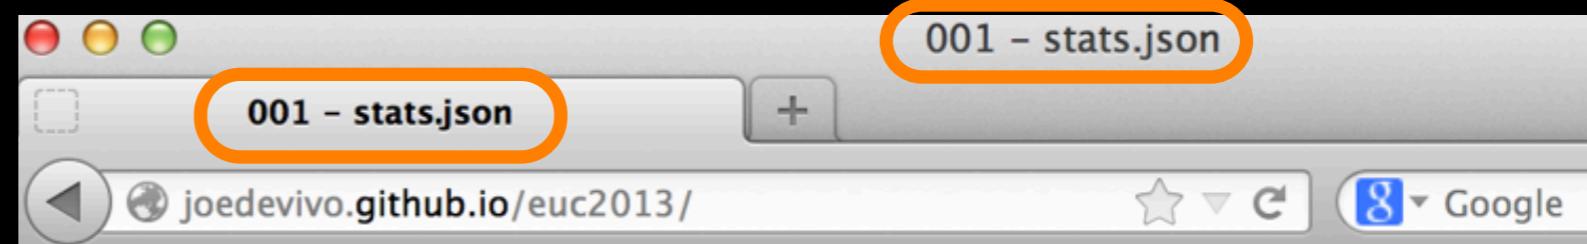
Use Thing

these → code: 001

# How to watch this talk

go here → <http://joedevivo.github.io/euc2013>

that! ➔



The screenshot shows a web browser window with the title bar "001 - stats.json". The address bar contains the URL "joedevivo.github.io/euc2013/". Below the address bar, there is a navigation bar with links for "Next" and "joedevivo/euc2013/snippets/001 - stats.json". The main content area displays a JSON object with 17 properties, all of which have a value of 0. The JSON object is as follows:

```
1 /* http://localhost:8091/stats riak ee 1.2.1 */
2 {
3     "vnode_gets":0,
4     "vnode_gets_total":0,
5     "vnode_puts":0,
6     "vnode_puts_total":0,
7     "vnode_index_reads":0,
8     "vnode_index_reads_total":0,
9     "vnode_index_writes":0,
10    "vnode_index_writes_total":0,
11    "vnode_index_writes_postings":0,
12    "vnode_index_writes_postings_total":0,
13    "vnode_index_deletes":0,
14    "vnode_index_deletes_total":0,
15    "vnode_index_deletes_postings":0,
16    "vnode_index_deletes_postings_total":0,
17    "node_gets":0.
```

Why are we here?

Add a feature to Riak

with riak\_test

using TDD

Add a feature to Riak  
with riak\_test  
using TDD

Riak is Basho's open-source,  
available, fault-tolerant,  
operationally simple,  
scalable, so you can rest  
databass

Add a feature to Riak

with riak\_test

using TDD

riak\_test is Basho's  
functional testing framework  
for testing distributed  
systems

# riak\_test

- does our release validation
- its roots are as a TDD framework

# Add a feature to Riak with riak\_test



using TDD

# Feature?

## JMX Monitoring for Riak

# Why JMX?

- It's pretty isolated
- We don't need to talk about complexities of riak
- It's closed source!
- I worked on it

Umm... hasn't Riak  
JMX been around for a  
while?

# What is Riak JMX?

/stats exposed!  
(in theory)

# what's /stats?

- <http://riak-node:8098/stats>
- riak-admin status
- general information and statistics about the node

code: 001

```
{  
    "vnode_gets":0,  
    "vnode_gets_total":0,  
    "vnode_puts":0,  
    "vnode_puts_total":0,  
    "vnode_index_reads":0,  
    "vnode_index_reads_total":0,  
    "vnode_index_writes":0,  
    "vnode_index_writes_total":0,  
    "vnode_index_writes_postings":0,  
    "vnode_index_writes_postings_total":0,  
    "vnode_index_deletes":0,  
    "vnode_index_deletes_total":0,  
    "vnode_index_deletes_postings":0,  
    "vnode_index_deletes_postings_total":0,  
    "node_gets":0,  
    "node_gets_total":0,  
    "node_get_fsm_siblings_mean":0,  
    "node_get_fsm_siblings_median":0,  
    "node_get_fsm_siblings_95":0,  
    "node_get_fsm_siblings_99":0,  
    "node_get_fsm_siblings_100":0,  
    "node_get_fsm_objsize_mean":0,  
    "node_get_fsm_objsize_median":0,  
    "node_get_fsm_objsize_95":0,  
    "node_get_fsm_objsize_99":0,  
    "node_get_fsm_objsize_100":0,  
}
```

```
[// → java -cp riak_jmx.jar com.basho.riak.jmx.Dump localhost 41110
{"CPUNProcs":145},
{"MemAllocated":5888280064},
 {"NodeGets":0},
 {"NodeGetsTotal":0},
 {"NodePuts":0},
 {"NodePutsTotal":0},
 {"MemTotal":7924740096},
 {"VnodeGets":0},
 {"VnodeGetsTotal":0},
 {"VnodePuts":0},
 {"VnodePutsTotal":0},
 {"PbcActive":0},
 {"PbcConnects":0},
 {"PbcConnectsTotal":0},
 {"NodeName":"dev1@127.0.0.1"},
 {"RingCreationSize":64},
 {"CpuAvg1":320},
 {"CpuAvg5":279},
 {"CpuAvg15":274},
 {"NodeGetFsmTimeMax":0},
 {"NodeGetFsmTimeMean":0},
 {"NodeGetFsmTimeMedian":0},
 {"NodeGetFsmTime95":0},
 {"NodeGetFsmTime99":0},
 {"NodePutFsmTimeMax":0},
 {"NodePutFsmTimeMean":0},
 {"NodePutFsmTimeMedian":0},
 {"NodePutFsmTime95":0},
 {"NodePutFsmTime99":0},
 {"ReadRepairs":0},
 {"ReadRepairsTotal":0}
]
code: 002
```

```
{
  "vnode_gets":0,
  "vnode_gets_total":0,
  "vnode_puts":0,
  "vnode_puts_total":0,
  "vnode_index_reads":0,
  "vnode_index_reads_total":0,
  "vnode_index_writes":0,
  "vnode_index_writes_total":0,
  "vnode_index_writes_postings":0,
  "vnode_index_writes_postings_total":0,
  "vnode_index_deletes":0,
  "vnode_index_deletes_total":0,
  "vnode_index_deletes_postings":0,
  "vnode_index_deletes_postings_total":0,
  "node_gets":0,
  "node_gets_total":0,
  "node_get_fsm_siblings_mean":0,
  "node_get_fsm_siblings_median":0,
  "node_get_fsm_siblings_95":0,
  "node_get_fsm_siblings_99":0,
  "node_get_fsm_siblings_100":0,
  "node_get_fsm_objsize_mean":0,
  "node_get_fsm_objsize_median":0,
  "node_get_fsm_objsize_95":0,
  "node_get_fsm_objsize_99":0,
  "node_get_fsm_objsize_100":0,
  "node_get_fsm_time_mean":0,
  "node_get_fsm_time_median":0,
  "node_get_fsm_time_95":0,
  "node_get_fsm_time_99":0,
  "node_get_fsm_time_100":0,
  "node_puts":0,
  "node_puts_total":0,
  "node_put_fsm_time_mean":0,
  "node_put_fsm_time_median":0,
  "node_put_fsm_time_95":0,
  "node_put_fsm_time_99":0,
  "node_put_fsm_time_100":0,
  "read_repairs":0,
  "read_repairs_total":0,
  "coord_redirs_total":0,
  "executing_mappers":0,
  "precommit_fail":0,
  "postcommit_fail":0,
  "pbc_active":0,
  "pbc_connects":0,
  "pbc_connects_total":0,
  "cpu_nprocs":155,
  "cpu_avg1":233,
  "cpu_avg5":228,
  "cpu_avg15":241,
  "mem_total":8164060000,
  "mem_allocated":7943324000,
  "nodename":"dev1@127.0.0.1",
  "connected_nodes":[],
  "sys_driver_version":"2.0",
  "sys_global_heaps_size":0,
  "sys_heap_type":"private",
  "sys_logical_processors":4,
  "sys_otp_release":"R15B01",
  "sys_process_count":1440,
  "sys_smp_support":true,
  "sys_system_version":"Erlang R15B01 (erts-5.9.1) [source] [64-bit] [smp:4:4] [async-threads:64] [kernel-poll:true]",
  "sys_system_architecture":"i386-apple-darwin12.2.1",
  "sys_threads_enabled":true,
  "sys_thread_pool_size":64,
  "sys_wordsize":8,
  "ring_members":["dev1@127.0.0.1"],
  "ring_num_partitions":64,
  "ring_ownership":["[dev1@127.0.0.1,64]"],
  "ring_creation_size":64,
  "storage_backend":"riak_kv_bitcask_backend",
  "erlydtl_version":"0.7.0",
  "riak_control_version":"1.2.0",
  "cluster_info_version":"1.2.2",
  "riak_jmx_version":"1.2.0",
  "riak_snmp_version":"1.2.0",
  "riak_repl_version":"1.2.1",
  "ranch_version":"0.2.1",
  "riak_api_version":"1.2.0",
  "riak_search_version":"1.2.1",
  "merge_index_version":"1.2.1",
  "riak_kv_version":"1.2.1",
  "riak_pipe_version":"1.2.1",
  "riak_core_version":"1.2.1",
  "lager_version": "1.2.0",
  "syntax_tools_version": "1.6.8",
  "compiler_version": "4.8.1",
  "bitcask_version": "1.5.2",
  "basho_stats_version": "1.0.2",
  "luke_version": "0.2.5",
  "webmachine_version": "1.9.2",
  "mochiweb_version": "1.5.1",
  "inets_version": "5.0",
  "erlang_js_version": "1.2.1",
  "mnesia_version": "4.7",
  "snmp_version": "4.22",
  "runtime_tools_version": "1.8.8",
  "os_mon_version": "2.2.9",
  "riak_sysmon_version": "1.1.2",
  "ssl_version": "5.0.1",
  "public_key_version": "0.15",
  "crypto_version": "2.1",
  "sasl_version": "2.2.1",
  "stdlib_version": "1.18.1",
  "kernel_version": "2.15.1",
  "memory_total": 30986792,
  "memory_processes": 9083110,
  "memory_processes_used": 9083050,
  "memory_system": 21903682,
  "memory_atom": 561761,
  "memory_atom_used": 546330,
  "memory_binary": 71424,
  "memory_code": 13914313,
  "memory_ets": 1977760,
  "ignored_gossip_total": 0,
  "rings_reconciled_total": 0,
  "rings_reconciled": 0,
  "gossip_received": 0,
  "rejected_handoffs": 0,
  "handoff_timeouts": 0,
  "converge_delay_min": 0,
  "converge_delay_max": 0,
  "converge_delay_mean": 0,
  "converge_delay_last": 0,
  "rebalance_delay_min": 0
}

[{"CPUNProcs":145,
  "MemAllocated":5888280064,
  {"NodeGets":0},
  {"NodeGetsTotal":0},
  {"NodePuts":0},
  {"NodePutsTotal":0},
  {"MemTotal":7924740096},
  {"VnodeGets":0},
  {"VnodeGetsTotal":0},
  {"VnodePuts":0},
  {"VnodePutsTotal":0},
  {"PbcActive":0},
  {"PbcConnects":0},
  {"PbcConnectsTotal":0},
  {"NodeName":"dev1@127.0.0.1"},
  {"RingCreationSize":64},
  {"CpuAvg1":320},
  {"CpuAvg5":279},
  {"CpuAvg15":274},
  {"NodeGetFsmTimeMax":0},
  {"NodeGetFsmTimeMean":0},
  {"NodeGetFsmTimeMedian":0},
  {"NodeGetFsmTime95":0},
  {"NodeGetFsmTime99":0},
  {"NodePutFsmTimeMax":0},
  {"NodePutFsmTimeMean":0},
  {"NodePutFsmTimeMedian":0},
  {"NodePutFsmTime95":0},
  {"NodePutFsmTime99":0},
  {"ReadRepairs":0},
  {"ReadRepairsTotal":0}
]
```

history lesson:

Riak JMX predates  
riak\_test

Why does that matter?

We can't go full big  
TDD on it

Let's go back in time to  
the riak 1.3 release

# Reuse!

# verify\_stats.erl

code: 003

```
confirm() ->  
    Nodes = [Node1] = rt:deploy_nodes(1),  
    ?assertEqual(ok,  
                rt:wait_until_nodes_ready(Nodes)),  
  
    Nodes = [Node1, Node2, Node3] = rt:deploy_nodes(3),
```

```
Stats1 = get_stats(Node1),  
  
get_stats(Node) ->  
    timer:sleep(10000),  
    StatString = os:cmd(  
        io_lib:format("curl -s -S ~s/stats",  
                      [rt:http_url(Node)])),  
    {struct, Stats} = mochijson2:decode(StatString),  
    Stats.
```

```
%% make sure a set of stats have valid values
verify_nz(Stats1, [<<"cpu_nprocs">>,
                  <<"mem_total">>,
                  <<"mem_allocated">>,
                  <<"sys_logical_processors">>,
                  <<"sys_process_count">>,
                  <<"sys_thread_pool_size">>,
                  <<"sys_wordsize">>,
                  <<"ring_num_partitions">>,
                  <<"ring_creation_size">>,
                  <<"memory_total">>,
                  <<"memory_processes">>,
                  <<"memory_processes_used">>,
                  <<"memory_system">>,
                  <<"memory_atom">>,
                  <<"memory_atom_used">>,
                  <<"memory_binary">>,
                  <<"memory_code">>,
                  <<"memory_ets">>]) ,
```

```
verify_nz(Props, Keys) ->
[?assertNotEqual(
    proplists:get_value(Key, Props, 0),
    0)
 || Key <- Keys].
```

```
%% make sure a set of stats have valid values
verify_nz(Stats1, [<<"cpu_nprocs">>,
                  <<"mem_total">>,
                  <<"mem_allocated">>,
                  <<"sys_logical_processors">>,
                  <<"sys_process_count">>,
                  <<"sys_thread_pool_size">>,
                  <<"sys_wordsize">>,
                  <<"ring_num_partitions">>,
                  <<"ring_creation_size">>,
                  <<"memory_total">>,
                  <<"memory_processes">>,
                  <<"memory_processes_used">>,
                  <<"memory_system">>,
                  <<"memory_atom">>,
                  <<"memory_atom_used">>,
                  <<"memory_binary">>,
                  <<"memory_code">>,
                  <<"memory_ets">>]) ,
```

```
lager:info(
    "perform 5 x  PUT and a GET to increment the stats"),
lager:info(
    "as the stat system only does calcs for > 5 readings"),

C = rt:httpc(Node1),

[rt:httpc_write(C,
                <<"systest">>,
                <<X>>,
                <<"12345">>)
 || X <- lists:seq(1, 5)],

[rt:httpc_read(C,
                <<"systest">>,
                <<X>>)
 || X <- lists:seq(1, 5)],
```

```
Stats1 = get_stats(Node1),
```

I'm still bound!  remember me?

```
Stats2 = get_stats(Node1),
```

```
%% make sure the stats that were supposed to increment did
verify_inc(Stats1, Stats2, [{<<"node_gets">>, 10},
                           {<<"node_puts">>, 5},
                           {<<"node_gets_total">>, 10},
                           {<<"node_puts_total">>, 5},
                           {<<"vnode_gets">>, 30},
                           {<<"vnode_puts">>, 15},
                           {<<"vnode_gets_total">>, 30},
                           {<<"vnode_puts_total">>, 15}]),
```

```
verify_inc(Stats1, Stats2, [{<<"node_gets">>, 10},  
 ...]),  
  
verify_inc(Prev, Props, Keys) ->  
[begin  
    Old = proplists:get_value(Key, Prev, 0),  
    New = proplists:get_value(Key, Props, 0),  
    lager:info("~s: ~p -> ~p (expected ~p)",  
              [Key, Old, New, Old + Inc]),  
    ?assertEqual(New, (Old + Inc))  
end || {Key, Inc} <- Keys].
```

```
%% make sure the stats that were supposed to increment did
verify_inc(Stats1, Stats2, [{<<"node_gets">>, 10},
                           {<<"node_puts">>, 5},
                           {<<"node_gets_total">>, 10},
                           {<<"node_puts_total">>, 5},
                           {<<"vnode_gets">>, 30},
                           {<<"vnode_puts">>, 15},
                           {<<"vnode_gets_total">>, 30},
                           {<<"vnode_puts_total">>, 15}]),
```

```
%% verify that fsm times were tallied
verify_nz(Stats2, [<<"node_get_fsm_time_mean">>,
                   <<"node_get_fsm_time_median">>,
                   <<"node_get_fsm_time_95">>,
                   <<"node_get_fsm_time_99">>,
                   <<"node_get_fsm_time_100">>,
                   <<"node_put_fsm_time_mean">>,
                   <<"node_put_fsm_time_median">>,
                   <<"node_put_fsm_time_95">>,
                   <<"node_put_fsm_time_99">>,
                   <<"node_put_fsm_time_100">>] ),
```





pass.

large font because this is great!

What's different about  
JMX?

Oh no!

# Dump.java

code: 004

```
for( ObjectInstance instance : beans )
{
    MBeanInfo info = mbs.getMBeanInfo( instance.getObjectName() );
    MBeanAttributeInfo[] attrs = info.getAttributes();
    for (MBeanAttributeInfo attr: attrs)
    {
        String name = attr.getName();
        Object value = mbs.getAttribute(riakBeanName, name);
        results.object().key(name);
        if(value instanceof Number)
        {
            results.value(((Number)value).longValue());
        } else {
            results.value(value);
        }
        results.endObject();
    }
}
System.out.println(results.endArray().toString());
```

# jmx\_verify.erl

code: 005

```
JMXDumpCmd = jmx_dump_cmd(IP, JMXPort),
JMX1 = jmx_dump(JMXDumpCmd),  
  
jmx_jar_path() ->
    %% Find riak_jmx.jar
    DepsPath = rt:get_deps(),
    Deps = string:tokens(os:cmd("ls " ++ DepsPath), "\n"),
    [RiakJMX] = lists:filter(fun(X) -
> string:str(X, "riak_jmx") == 1 end, Deps),
    filename:join([DepsPath, RiakJMX, "priv", "riak_jmx.jar"]).
  
jmx_dump_cmd(IP, Port) ->
    io_lib:format("java -cp ~s com.basho.riak.jmx.Dump ~s ~p",
        [jmx_jar_path(), IP, Port]).  
  
jmx_dump(Cmd) ->
    timer:sleep(40000), %% JMX only updates every 30seconds
    Output = string:strip(os:cmd(Cmd), both, $\n),
    JSONOutput = mochijson2:decode(Output),
    [ {Key, Value} || {struct, [{Key, Value}]} ] <- JSONOutput.
```



# How do we run it?

```
→ riak_ee git:(master) make devrel
...
→ riak_ee git:(master) ./riak_test/bin/rteedev-current.sh
Making /Users/joe/dev/basho/riak_ee the current release:
- Determining version: riak_ee-1.4wip-70-gf5768c5-master
- Resetting existing /Users/joe/riak_test_ee
- Removing and recreating /Users/joe/riak_test_ee/current
- Copying devrel to /Users/joe/riak_test_ee/current
- Writing /Users/joe/riak_test_ee/current/VERSION
- Reinitializing git state
→ riak_ee git:(master) cd ../riak_test
→ riak_test git:(master) make clean && make
→ riak_test git:(master) ./riak_test -c ee -t jmx_verify
```

A close-up photograph of a large pile of wrapped hard candies. The candies are wrapped in clear plastic with colorful foil liners. Visible flavor labels include "BLUE RASPBERRY", "Lemon", and "GRAPE". The candies are in various shapes, some being round and others more irregular. They are piled high, filling the frame.

The test should work!

It doesn't

Parity

We'll rage on this later

code: 006

```
===== jmx_verify failure stack trace =====
{{assertNotEqual_failed,[{module,jmx_verify},
    {line,113},
    {expression,"{ Key , 0 }"},
    {value,{<<"cpu_nprocs">>,0}}]},,
[{jmx_verify,'-verify_nz/2-fun-0-',2,
    [{file,"tests/jmx_verify.erl"},{line,113}]}],
{jmx_verify,'-verify_nz/2-lc$^0/1-0-',2,
    [{file,"tests/jmx_verify.erl"},{line,113}]},
{jmx_verify,confirm,0,
    [{file,"tests/jmx_verify.erl"},{line,42}]}}}
```

---



**THERE ARE NO `cpu_nprocs`!**

cpu_nprocs	CPUNProcs
mem_allocated	MemAllocated
mem_total	MemTotal
node_gets	NodeGets
node_gets_total	NodeGetsTotal
node_puts	NodePuts
node_puts_total	NodePutsTotal
vnode_gets	VnodeGets
vnode_gets_total	VnodeGetsTotal
vnode_puts	VnodePuts
vnode_puts_total	VnodePutsTotal
pbc_active	PbcActive
pbc_connects	PbcConnects
pbc_connects_total	PbcConnectsTotal
nodename	NodeName
ring_creation_size	RingCreationSize
cpu_avg1	CpuAvg1
cpu_avg5	CpuAvg5
cpu_avg15	CpuAvg15
node_get_fsm_time_95	NodeGetFsmTime95
node_get_fsm_time_99	NodeGetFsmTime99
node_get_fsm_time_100	NodeGetFsmTimeMax
node_get_fsm_time_mean	NodeGetFsmTimeMean
node_get_fsm_time_median	NodeGetFsmTimeMedian
node_put_fsm_time_95	NodePutFsmTime95
node_put_fsm_time_99	NodePutFsmTime99
node_put_fsm_time_100	NodePutFsmTimeMax
node_put_fsm_time_mean	NodePutFsmTimeMean
node_put_fsm_time_median	NodePutFsmTimeMedian
read_repairs	ReadRepairs
read_repairs_total	ReadRepairsTotal

code: 007

code: 007

CPUNProcs → cpu\_nprocs

```
process_key(<<"CPUNProcs">>) -> <<"cpu_nprocs">>;
process_key(<<"MemAllocated">>) -> <<"mem_allocated">>;
process_key(<<"MemTotal">>) -> <<"mem_total">>;
process_key(<<"NodeGets">>) -> <<"node_gets">>;
process_key(<<"NodeGetsTotal">>) -> <<"node_gets_total">>;
process_key(<<"NodePuts">>) -> <<"node_puts">>;
process_key(<<"NodePutsTotal">>) -> <<"node_puts_total">>;
process_key(<<"VnodeGets">>) -> <<"vnode_gets">>;
process_key(<<"VnodeGetsTotal">>) -> <<"vnode_gets_total">>;
process_key(<<"VnodePuts">>) -> <<"vnode_puts">>;
process_key(<<"VnodePutsTotal">>) -> <<"vnode_puts_total">>;
process_key(<<"PbcActive">>) -> <<"pbc_active">>;
process_key(<<"PbcConnects">>) -> <<"pbc_connects">>;
process_key(<<"PbcConnectsTotal">>) -> <<"pbc_connects_total">>;
process_key(<<"NodeName">>) -> <<"nodename">>;
process_key(<<"RingCreationSize">>) -> <<"ring_creation_size">>;
process_key(<<"CpuAvg1">>) -> <<"cpu_avg1">>;
process_key(<<"CpuAvg5">>) -> <<"cpu_avg5">>;
process_key(<<"CpuAvg15">>) -> <<"cpu_avg15">>;
process_key(<<"NodeGetFsmTime95">>) -> <<"node_get_fsm_time_95">>;
process_key(<<"NodeGetFsmTime99">>) -> <<"node_get_fsm_time_99">>;
process_key(<<"NodeGetFsmTimeMax">>) -> <<"node_get_fsm_time_100">>;
process_key(<<"NodeGetFsmTimeMean">>) -> <<"node_get_fsm_time_mean">>;
process_key(<<"NodeGetFsmTimeMedian">>) -> <<"node_get_fsm_time_median">>;
process_key(<<"NodePutFsmTime95">>) -> <<"node_put_fsm_time_95">>;
process_key(<<"NodePutFsmTime99">>) -> <<"node_put_fsm_time_99">>;
process_key(<<"NodePutFsmTimeMax">>) -> <<"node_put_fsm_time_100">>;
process_key(<<"NodePutFsmTimeMean">>) -> <<"node_put_fsm_time_mean">>;
process_key(<<"NodePutFsmTimeMedian">>) -> <<"node_put_fsm_time_median">>;
process_key(<<"ReadRepairs">>) -> <<"read_repairs">>;
process_key(<<"ReadRepairsTotal">>) -> <<"read_repairs_total">>.
```

code: 008

```
09:32:46.305 [info] Wait until 'dev1@127.0.0.1' is pingable code: 009
09:32:46.306 [info] Check 'dev1@127.0.0.1' is a singleton
09:32:46.351 [info] Deployed nodes: ['dev1@127.0.0.1']
09:32:46.351 [info] Waiting for services [riak_kv] to start on
['dev1@127.0.0.1'].
09:32:47.354 [info] Wait until nodes are ready : ['dev1@127.0.0.1']
09:33:27.738 [info] perform 5 x PUT and a GET to increment the stats
09:33:27.738 [info] as the stat system only does calcs for > 5 readings
09:34:08.247 [info] node_gets: 0 -> 10 (expected 10)
09:34:08.247 [info] node_puts: 0 -> 5 (expected 5)
09:34:08.247 [info] node_gets_total: 0 -> 10 (expected 10)
09:34:08.248 [info] node_puts_total: 0 -> 5 (expected 5)
09:34:08.248 [info] vnode_gets: 0 -> 30 (expected 30)
09:34:08.248 [info] vnode_puts: 0 -> 15 (expected 15)
09:34:08.248 [info] vnode_gets_total: 0 -> 30 (expected 30)
09:34:08.248 [info] vnode_puts_total: 0 -> 15 (expected 15)
09:34:08.249 [info] Make PBC Connection
09:34:48.673 [info] pbc_connects_total: 0 -> 1 (expected 1)
09:34:48.673 [info] pbc_connects: 0 -> 1 (expected 1)
09:34:48.673 [info] pbc_active: 0 -> 1 (expected 1)
09:34:48.673 [info] Force Read Repair
09:35:29.144 [info] read_repairs_total: 0 -> 0 (expected 0)
09:35:29.145 [info] read_repairs: 0 -> 0 (expected 0)
09:36:09.521 [info] read_repairs_total: 0 -> 1 (expected 1)
09:36:09.521 [info] read_repairs: 0 -> 1 (expected 1)
09:36:09.522 [notice] jmx_verify Test Run Complete
```

Test Results:

jmx\_verify-bitcask: pass

Time passes...  
as time does

Somebody (me) wants  
to add features

add maven to the build

it should still work

```
10:47:31.088 [info] Wait until 'dev1@127.0.0.1' is pingable code: 010
10:47:31.089 [info] Check 'dev1@127.0.0.1' is a singleton
10:47:31.129 [info] Deployed nodes: ['dev1@127.0.0.1']
10:47:31.129 [info] Waiting for services [riak_kv] to start on
['dev1@127.0.0.1'].
10:47:31.131 [info] Wait until nodes are ready : ['dev1@127.0.0.1']
10:48:11.749 [info] perform 5 x PUT and a GET to increment the stats
10:48:11.749 [info] as the stat system only does calcs for > 5 readings
10:48:52.398 [info] node_gets: 0 -> 10 (expected 10)
10:48:52.398 [info] node_puts: 0 -> 5 (expected 5)
10:48:52.399 [info] node_gets_total: 0 -> 10 (expected 10)
10:48:52.399 [info] node_puts_total: 0 -> 5 (expected 5)
10:48:52.399 [info] vnode_gets: 0 -> 30 (expected 30)
10:48:52.399 [info] vnode_puts: 0 -> 15 (expected 15)
10:48:52.399 [info] vnode_gets_total: 0 -> 30 (expected 30)
10:48:52.400 [info] vnode_puts_total: 0 -> 15 (expected 15)
10:48:52.400 [info] Make PBC Connection
10:49:32.908 [info] pbc_connects_total: 0 -> 1 (expected 1)
10:49:32.909 [info] pbc_connects: 0 -> 1 (expected 1)
10:49:32.909 [info] pbc_active: 0 -> 1 (expected 1)
10:49:32.909 [info] Force Read Repair
10:50:13.448 [info] read_repairs_total: 0 -> 0 (expected 0)
10:50:13.449 [info] read_repairs: 0 -> 0 (expected 0)
10:50:53.829 [info] read_repairs_total: 0 -> 1 (expected 1)
10:50:53.829 [info] read_repairs: 0 -> 1 (expected 1)
10:50:53.829 [notice] jmx_verify Test Run Complete
```

Test Results:

jmx\_verify-bitcask: pass

A nasty rumor

We're going to need to  
kill some processes

So what do those  
processes look like?

# Add this!

```
?assert(fail),
```

# Now run it!

```
→ make clean && make  
→ ./riak_test -c ee -t jmx_verify
```

...

So, we find ourselves in a tricky situation here.

You've run a single test, and it has failed.  
Would you like to leave Riak running in order  
to debug?

[Y/n] Y



# What's it look like?

```
→ ps -ax
96410 ??          0:00.57 /usr/bin/java -server -
Dcom.sun.management.jmxremote.authenticate=false -
Dcom.sun.management.jmxremote.ssl=false -
Dcom.sun.management.jmxremote.port=41111
→ kill 96410
→ ps -ax
97142 ??          0:00.43 /usr/bin/java -server -
Dcom.sun.management.jmxremote.authenticate=false -
Dcom.sun.management.jmxremote.ssl=false -
Dcom.sun.management.jmxremote.port=41111
```

Let's kill them all!

```
kill_jmx() ->
{0, JMXProcs} = rt:cmd(
  "ps -ax | grep com.sun.management.jmxremote"),
Pids = filter_procs(string:tokens(JMXProcs, "\n")),
[ rt:cmd("kill -9 " ++ Pid) || Pid <- Pids].
```

```
filter_procs([]) -> [];
filter_procs([Proc|T]) ->
  case string:str(Proc, "java -server") of
    0 -> filter_procs(T);
    _ -> [hd(string:tokens(Proc, "\s"))
          | filter_procs(T)]
end.
```

```
test_supervision() ->
    JMXPort = 41111,
    Config = [{riak_jmx, [{enabled, true}, {port, JMXPort}] }],
    [Node] = rt:deploy_nodes(1, Config),

    [ kill_jmx() || _X <- lists:seq(0, 9)],

    case net_adm:ping(Node) of
        pang ->
            ?
    assertEqual("riak_jmx crash able to crash riak node", true);
        _ ->
            yay
    end,
ok.
```

It's too slow!

Let's listen on port 80

```
test_supervision() ->
    JMXPort = 80,
    Config = [{riak_jmx, [{enabled, true}, {port, JMXPort}] }],
    [Node] = rt:deploy_nodes(1, Config),
    timer:sleep(20000),
    case net_adm:ping(Node) of
        pang ->
            lager:error(
                "riak_jmx crash able to crash riak node"),
            ?assertEqual(
                "riak_jmx crash able to crash riak node", true);
        _ ->
            yay
    end,
```

ok.

# That triggers it!

```
===== jmx_verify failure stack trace =====
{{assertEqual_failed,
 [{module,jmx_verify},
  {line,176},
  {expression,"true"},
  {expected,"riak_jmx crash able to crash riak node"},
  {value,true}]},
 [{jmx_verify,'-test_supervision/0-fun-0-',1,
    [{file,"tests/jmx_verify.erl"},{line,176}]},
  {jmx_verify,test_supervision,0,[{file,"tests/
jmx_verify.erl"},{line,176}]},
  {jmx_verify,confirm,0,[{file,"tests/jmx_verify.erl"},{line,
29}]}]}
=====
```

# What should it do?

- Let the java process crash
- If it does, try 10 times like it used to
- If that fails, sleep for 10 minutes and try again

Let's add to the test to  
check for that non-  
existent functionality

```
%% Let's make sure the thing's restarting as planned
lager:info("calling riak_jmx:stop() to reset retry counters"), code: 016
rpc:call(Node, riak_jmx, stop, ["stopping for test purposes"]),

lager:info("loading lager backend on node"),
rt:load_modules_on_nodes([riak_test_lager_backend], [Node]),
ok = rpc:call(Node, gen_event, add_handler, [lager_event, riak_test_lager_backend, [info, false]]),
ok = rpc:call(Node, lager, set_loglevel, [riak_test_lager_backend, info]),

lager:info("Now we're capturing logs on the node, let's start jmx"),
lager:info("calling riak_jmx:start() to get these retries started"),
rpc:call(Node, riak_jmx, start, []),

timer:sleep(40000), %% wait 2000 millis per restart + fudge factor
Logs = rpc:call(Node, riak_test_lager_backend, get_logs, []),

lager:info("It can fail, it can fail 10 times"),
RetryCount = lists:foldl(
    fun(Log, Sum) -
        > case re:run(element(7, element(2, Log)), "JMX server monitor .* exited with
code .*\\. Retrying", []) of
            {match, _} -> 1 + Sum;
            _ -> Sum
        end
    end,
    0, Logs),
?assertEqual({retry_count, 10}, {retry_count, RetryCount}),
rt:stop(Node),
ok_ok.
```

```
===== jmx_verify failure stack trace =====
{function_clause,
 [{lists,foldl,
  [#Fun<jmx_verify.3.111466063>,0,{badrpc,nodedown}],
   [{file,"lists.erl"},{line,1196}]},
 {jmx_verify,test_supervision,0,
  [{file,"tests/jmx_verify.erl"},{line,198}]},
 {jmx_verify,confirm,0,
  [{file,"tests/jmx_verify.erl"},{line,29}]}]}
=====
```

code: 018

```
rpc:call(Node, riak_jmx, stop, ["stopping for test purposes"] ),  
  
stop(Reason) ->  
    error_logger:info_msg(io_lib:format("~p~n", [Reason])),  
    init:stop().
```

code: 018

```
rpc:call(Node, riak_jmx, stop, ["stopping for test purposes"] ),
```

```
stop(Reason) ->  
    error_logger:info_msg(io_lib:format("~p~n", [Reason])),  
    application:stop(riak_jmx).
```

```
===== jmx_verify failure stack trace =====
{{assertEqual_failed,
  [{module,jmx_verify},
   {line,205},
   {expression,"{ retry_count , RetryCount }"},  

   {expected,{retry_count,10}},
   {value,{retry_count,15}}]},  

 [{jmx_verify,'-test_supervision/0-fun-2-',2,
   [{file,"tests/jmx_verify.erl"},{line,205}]},
  {jmx_verify,test_supervision,0,[{file,"tests/jmx_verify.erl"},
  {line,205}]},
  {jmx_verify,confirm,0,[{file,"tests/jmx_verify.erl"},{line,
  29}]}]}
=====
```

# log the retry count

code: 020,021

```
lager:info(  
    "JMX server monitor ~s exited with code ~p. Retrying.",  
    [State#state.pid, Rc]),
```

code: 020,021

# log the retry count

lager:info(

"JMX server monitor ~s exited with code ~p. Retry #~p.",  
[State#state.pid, Rc, Retry]),

{expression,"{ retry\_count , RetryCount }"},  
{expected,{retry\_count,10}},  
{value,{retry\_count,15}}]},

13:23:26.395	[info]	JMX	...	40493	exited with code	1.	Retry #0.
13:23:29.081	[info]	JMX	...	40519	exited with code	1.	Retry #0.
13:23:31.756	[info]	JMX	...	40556	exited with code	1.	Retry #0.
13:23:34.451	[info]	JMX	...	40582	exited with code	1.	Retry #0.
13:23:37.124	[info]	JMX	...	40619	exited with code	1.	Retry #0.
13:23:39.867	[info]	JMX	...	40646	exited with code	1.	Retry #0.
13:23:42.529	[info]	JMX	...	40684	exited with code	1.	Retry #0.
13:23:45.174	[info]	JMX	...	40710	exited with code	1.	Retry #0.
13:23:47.824	[info]	JMX	...	40747	exited with code	1.	Retry #0.
13:23:50.460	[info]	JMX	...	40774	exited with code	1.	Retry #0.
13:23:53.153	[info]	JMX	...	40800	exited with code	1.	Retry #0.
13:23:55.862	[info]	JMX	...	40827	exited with code	1.	Retry #0.

OOPS! Forgot to  
increment the retry  
counter

code: 022

# log the retry count

```
lager:info(  
    "JMX server monitor ~s exited with code ~p. Retry #~p.",  
    [State#state.pid, Rc, Retry]),
```

14:22:52.301 [info] "stopping for test purposes"	
14:22:53.004 [info] JMX ... 47923 exited with code 1.	Retry #0.
14:22:55.636 [info] JMX ... 47960 exited with code 1.	Retry #1.
14:22:58.305 [info] JMX ... 47985 exited with code 1.	Retry #2.
14:23:00.979 [info] JMX ... 48012 exited with code 1.	Retry #3.
14:23:03.640 [info] JMX ... 48049 exited with code 1.	Retry #4.
14:23:06.428 [info] JMX ... 48074 exited with code 1.	Retry #5.
14:23:09.076 [info] JMX ... 48111 exited with code 1.	Retry #6.
14:23:11.795 [info] JMX ... 48138 exited with code 1.	Retry #7.
14:23:14.484 [info] JMX ... 48175 exited with code 1.	Retry #8.
14:23:17.214 [info] JMX ... 48200 exited with code 1.	Retry #9.

Cool!

Now, let's bring parity

cpu_nprocs	CPUNProcs
mem_allocated	MemAllocated
mem_total	MemTotal
node_gets	NodeGets
node_gets_total	NodeGetsTotal
node_puts	NodePuts
node_puts_total	NodePutsTotal
vnode_gets	VnodeGets
vnode_gets_total	VnodeGetsTotal
vnode_puts	VnodePuts
vnode_puts_total	VnodePutsTotal
pbc_active	PbcActive
pbc_connects	PbcConnects
pbc_connects_total	PbcConnectsTotal
nodename	NodeName
ring_creation_size	RingCreationSize
cpu_avg1	CpuAvg1
cpu_avg5	CpuAvg5
cpu_avg15	CpuAvg15
node_get_fsm_time_95	NodeGetFsmTime95
node_get_fsm_time_99	NodeGetFsmTime99
node_get_fsm_time_100	NodeGetFsmTimeMax
node_get_fsm_time_mean	NodeGetFsmTimeMean
node_get_fsm_time_median	NodeGetFsmTimeMedian
node_put_fsm_time_95	NodePutFsmTime95
node_put_fsm_time_99	NodePutFsmTime99
node_put_fsm_time_100	NodePutFsmTimeMax
node_put_fsm_time_mean	NodePutFsmTimeMean
node_put_fsm_time_median	NodePutFsmTimeMedian
read_repairs	ReadRepairs
read_repairs_total	ReadRepairsTotal

**CPUNProcs** → `cpu_nprocs`

code: 023

```
process_key(<<"CPUNProcs">>) -> <<"cpu_nprocs">>;  
process_key(<<"MemAllocated">>) -> <<"mem_allocated">>;  
process_key(<<"MemTotal">>) -> <<"mem_total">>;  
process_key(<<"NodeGets">>) -> <<"node_gets">>;  
process_key(<<"NodeGetsTotal">>) -> <<"node_gets_total">>;  
process_key(<<"NodePuts">>) -> <<"node_puts">>;  
process_key(<<"NodePutsTotal">>) -> <<"node_puts_total">>;  
process_key(<<"VnodeGets">>) -> <<"vnode_gets">>;  
process_key(<<"VnodeGetsTotal">>) -> <<"vnode_gets_total">>;  
process_key(<<"VnodePut">>) -> <<"vnode_puts">>;  
process_key(<<"VnodePutTotal">>) -> <<"vnode_puts_total">>;  
process_key(<<"PbcActive">>) -> <<"pbc_active">>;  
process_key(<<"PbcConnects">>) -> <<"pbc_connects">>;  
process_key(<<"PbcConnectsTotal">>) -> <<"pbc_connects_total">>;  
process_key(<<"NodeName">>) -> <<"node_name">>;  
process_key(<<"RingCreationSize">>) -> <<"ring_creation_size">>;  
process_key(<<"CpuAvg1">>) -> <<"cpu_avg1">>;  
process_key(<<"CpuAvg5">>) -> <<"cpu_avg5">>;  
process_key(<<"CpuAvg15">>) -> <<"cpu_avg15">>;  
process_key(<<"NodeGetFsmTime95">>) -> <<"node_get_fsm_time_95">>;  
process_key(<<"NodeGetFsmTime99">>) -> <<"node_get_fsm_time_99">>;  
process_key(<<"NodeGetFsmTimeMax">>) -> <<"node_get_fsm_time_100">>;  
process_key(<<"NodeGetFsmTimeMean">>) -> <<"node_get_fsm_time_mean">>;  
process_key(<<"NodeGetFsmTimeMedian">>) -> <<"node_get_fsm_time_median">>;  
process_key(<<"NodePutFsmTime95">>) -> <<"node_put_fsm_time_95">>;  
process_key(<<"NodePutFsmTime99">>) -> <<"node_put_fsm_time_99">>;  
process_key(<<"NodePutFsmTimeMax">>) -> <<"node_put_fsm_time_100">>;  
process_key(<<"NodePutFsmTimeMean">>) -> <<"node_put_fsm_time_mean">>;  
process_key(<<"NodePutFsmTimeMedian">>) -> <<"node_put_fsm_time_median">>;  
process_key(<<"ReadRepairs">>) -> <<"read_repairs">>;  
process_key(<<"ReadRepairsTotal">>) -> <<"read_repairs_total">>.
```



```
===== jmx_verify failure stack trace =====
{{assertNotEqual_failed,[{module,jmx_verify},
    {line,115},
    {expression,"{ Key , 0 }"},{value,{<<"cpu_nprocs">>,0}}]},,
[{jmx_verify,'-verify_nz/2-fun-0-',2,
    [{file,"tests/jmx_verify.erl"},{line,115}]},,
{jmx_verify,'-verify_nz/2-lc$^0/1-0-',2,
    [{file,"tests/jmx_verify.erl"},{line,115}]},,
{jmx_verify,confirm,0,[{file,"tests/jmx_verify.erl"},{line,
44}]}]}
```

---

# Handwavey Java

- class & interface
- javassist
- curl vs jinterface
- crashing
- permgen space

11:02:52.775 [info] Wait until 'dev1@127.0.0.1' is pingable code: 025  
11:02:52.775 [info] Wait until the ring manager is up on 'dev1@127.0.0.1'  
11:02:52.776 [info] Check 'dev1@127.0.0.1' is a singleton  
11:02:52.776 [info] Deployed nodes: ['dev1@127.0.0.1']  
11:02:52.776 [info] Waiting for services [riak\_kv] to start on  
['dev1@127.0.0.1'].  
11:03:01.805 [info] Wait until nodes are ready : ['dev1@127.0.0.1']  
11:03:42.309 [info] perform 5 x PUT and a GET to increment the stats  
11:03:42.310 [info] as the stat system only does calcs for > 5 readings  
11:04:22.968 [info] node\_gets: 0 -> 10 (expected 10)  
11:04:22.969 [info] node\_puts: 0 -> 5 (expected 5)  
11:04:22.969 [info] node\_gets\_total: 0 -> 10 (expected 10)  
11:04:22.969 [info] node\_puts\_total: 0 -> 5 (expected 5)  
11:04:22.969 [info] vnode\_gets: 0 -> 30 (expected 30)  
11:04:22.969 [info] vnode\_puts: 0 -> 15 (expected 15)  
11:04:22.970 [info] vnode\_gets\_total: 0 -> 30 (expected 30)  
11:04:22.970 [info] vnode\_puts\_total: 0 -> 15 (expected 15)  
11:04:22.970 [info] Make PBC Connection  
11:05:03.469 [info] pbc\_connects\_total: 0 -> 1 (expected 1)  
11:05:03.469 [info] pbc\_connects: 0 -> 1 (expected 1)  
11:05:03.470 [info] pbc\_active: 0 -> 1 (expected 1)  
11:05:03.470 [info] Force Read Repair  
11:05:44.104 [info] read\_repairs\_total: 0 -> 0 (expected 0)  
11:05:44.104 [info] read\_repairs: 0 -> 0 (expected 0)  
11:06:24.543 [info] read\_repairs\_total: 0 -> 1 (expected 1)  
11:06:24.543 [info] read\_repairs: 0 -> 1 (expected 1)  
11:06:24.544 [notice] jmx\_verify Test Run Complete

Test Results:

jmx\_verify-bitcask: pass

We did wrong

DynamicMBean is the  
HashMap of JMX Beans

09:42:05.367 [info] Wait until 'dev1@127.0.0.1' is pingable code: 026  
09:42:05.368 [info] Wait until the ring manager is up on 'dev1@127.0.0.1'  
09:42:05.369 [info] Check 'dev1@127.0.0.1' is a singleton  
09:42:05.370 [info] Deployed nodes: ['dev1@127.0.0.1']  
09:42:05.370 [info] Waiting for services [riak\_kv] to start on  
['dev1@127.0.0.1'].  
09:42:14.393 [info] Wait until nodes are ready : ['dev1@127.0.0.1']  
09:42:54.873 [info] perform 5 x PUT and a GET to increment the stats  
09:42:54.873 [info] as the stat system only does calcs for > 5 readings  
09:43:35.560 [info] node\_gets: 0 -> 10 (expected 10)  
09:43:35.560 [info] node\_puts: 0 -> 5 (expected 5)  
09:43:35.560 [info] node\_gets\_total: 0 -> 10 (expected 10)  
09:43:35.561 [info] node\_puts\_total: 0 -> 5 (expected 5)  
09:43:35.561 [info] vnode\_gets: 0 -> 30 (expected 30)  
09:43:35.561 [info] vnode\_puts: 0 -> 15 (expected 15)  
09:43:35.562 [info] vnode\_gets\_total: 0 -> 30 (expected 30)  
09:43:35.562 [info] vnode\_puts\_total: 0 -> 15 (expected 15)  
09:43:35.562 [info] Make PBC Connection  
09:44:16.102 [info] pbc\_connects\_total: 0 -> 1 (expected 1)  
09:44:16.103 [info] pbc\_connects: 0 -> 1 (expected 1)  
09:44:16.103 [info] pbc\_active: 0 -> 1 (expected 1)  
09:44:16.103 [info] Force Read Repair  
09:44:56.688 [info] read\_repairs\_total: 0 -> 0 (expected 0)  
09:44:56.688 [info] read\_repairs: 0 -> 0 (expected 0)  
09:45:37.242 [info] read\_repairs\_total: 0 -> 1 (expected 1)  
09:45:37.242 [info] read\_repairs: 0 -> 1 (expected 1)  
09:45:37.242 [notice] jmx\_verify Test Run Complete

Test Results:

jmx\_verify-bitcask: pass

# Jump to Conclusions

???

JUMP  
AGAIN

STRIKE  
OUT

COULD  
BE

LOSE  
ONE  
TURN

YES!

NO!

ACCEPT  
IT

GO  
WILD

ONE  
STEP  
BACK

THINK  
AGAIN

MOOT!

# Lessons Learned

TDD is like going to the gym

# Lessons Learned

riak\_test is better  
at CI than TDD

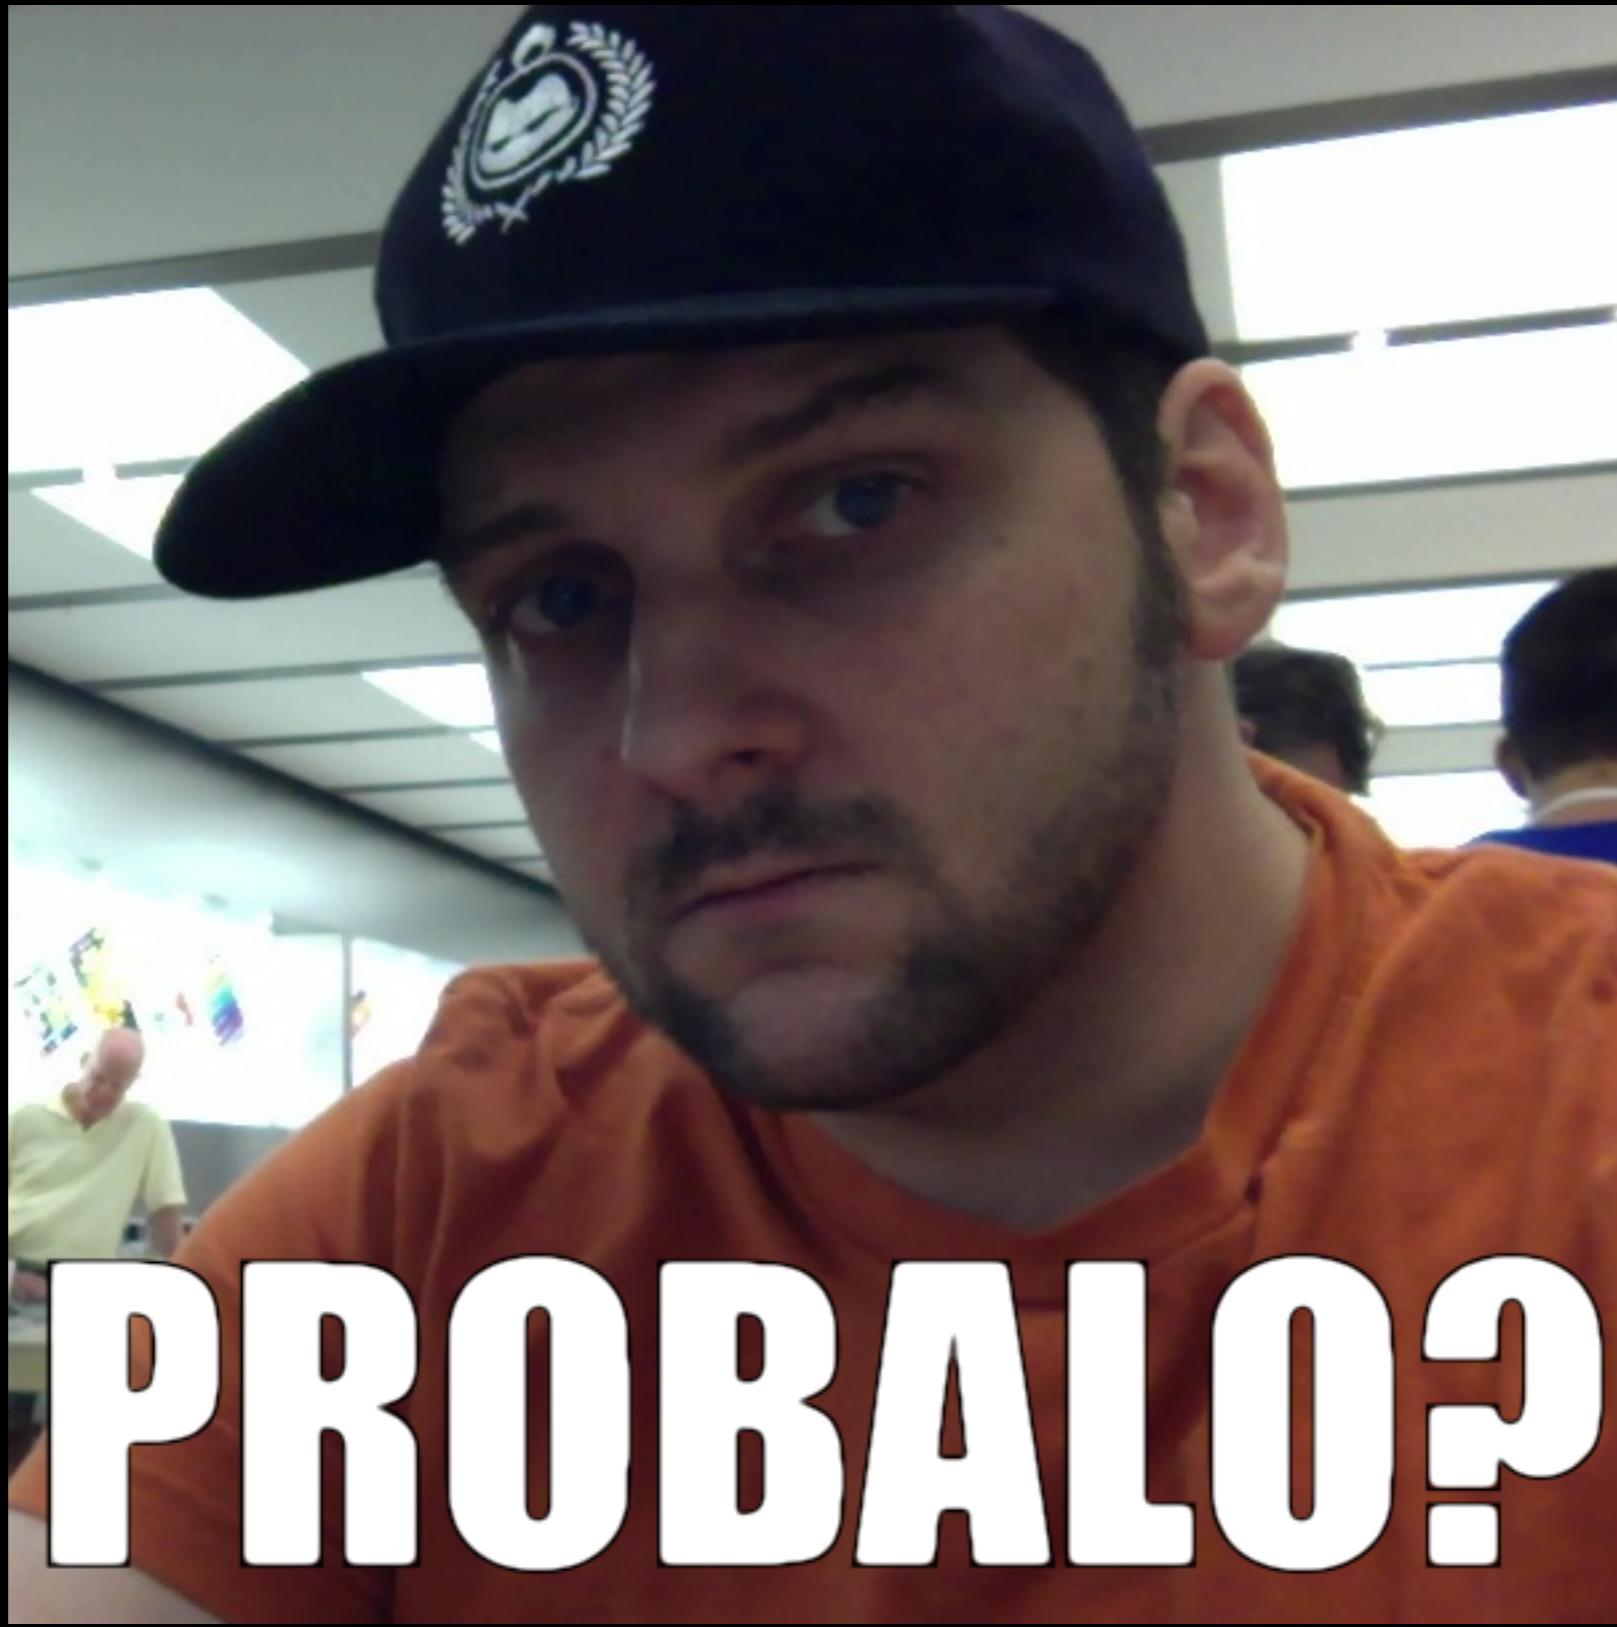
# Lessons Learned

It would have been hard  
to test the restart stuff  
manually

# find out more

<http://joedevivo.github.io/euc2013>

[http://github.com/basho/riak\\_test](http://github.com/basho/riak_test)



**PROBALO?**

riak\_test's greatest  
hits!

# basic\_command\_line

```
%% Stop node, to test console working
rt:console(Node, [
{expect, "\(\abort with ^G\)"}},
{send, "riak_core_ring_manager:get_my_ring()."},
{expect, "dict,"},
{send, "q()." },
{expect, "ok"}]),
rt:wait_until_unpingable(Node),
ok.
```

# counter converge

```
[N1, N2, N3, N4]=Nodes = rt:build_cluster(4),  
  
PartInfo = rt:partition([N1, N2], [N3, N4]),  
  
%% increment one side  
increment_counter(HP1, Key, 5),  
  
%% check value on one side is different from other  
[?assertEqual(13, get_counter(HP, Key)) || HP <- [HP1, HP2]],  
[?assertEqual(8, get_counter(HP, Key)) || HP <- [HP3, HP4]],  
  
%% decrement other side  
decrement_counter(HP3, Key, 2),  
  
%% verify values differ  
[?assertEqual(13, get_counter(HP, Key)) || HP <- [HP1, HP2]],  
[?assertEqual(6, get_counter(HP, Key)) || HP <- [HP3, HP4]],  
  
%% heal  
lager:info("Heal and check merged values"),  
ok = rt:heal(PartInfo),  
ok = rt:wait_for_cluster_service(Nodes, riak_kv),  
  
%% verify all nodes agree  
[?assertEqual(ok, rt:wait_until(HP, fun(N) ->  
                                         11 == get_counter(N, Key)  
                                       end)) || HP <- Hosts],  
  
pass.
```

# loaded upgrade

```
[begin
    exit(Sup, normal),
    lager:info("Upgrading ~p", [Node]),
    rt:upgrade(Node, current),
    {ok, NewSup} = rt_worker_sup:start_link([
        {concurrent, Concurrent},
        {node, Node},
        {backend, Backend},
        {version, current}
    ]),
    _NodeMon = init_node_monitor(Node, NewSup, self()),
    upgrade_recv_loop()

end || {{ok, Sup}, Node} <- Sups],
pass.
```

# loaded upgrade

